

Deep Learning – Assignment 1: Single-Word Audio Classification (FNN)

Library

In [1]: `!nvidia-smi`

Fri May 2 22:28:38 2025

NVIDIA-SMI 572.42				Driver Version: 572.42				CUDA Version: 12.8			
GPU		Name		Driver-Model		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan	Temp	Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M. MIG M.	
=====											
0	NVIDIA	GeForce	RTX 3060 ...	WDDM	00000000:01:00.0		Off	N/A			
N/A	51C	P8	13W / 75W	281MiB / 6144MiB		0%		Default N/A			

Processes:											
GPU	GI	CI	PID	Type	Process name				GPU Memory		
		ID	ID	Usage							
=====											
0	N/A	N/A	28676	C	...s\Python\Python310\python.exe				N/A		

```
In [2]: import os

import time
import torch
import torch.nn as nn
import torch.optim as optim
import torchaudio
import librosa
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, accuracy_score
from torch.utils.data import Dataset, DataLoader
from tqdm import tqdm
```

```
In [3]: # Check if GPU is available
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)
```

cuda

Extract Audio

```
In [4]: class AudioDataset(Dataset):
    def __init__(self, root_dir):
        self.root_dir = root_dir
        self.samples = []
        self.labels = []
        self.label_map = {}

        for idx, label in enumerate(sorted(os.listdir(root_dir))):
            label_path = os.path.join(root_dir, label)
            if os.path.isdir(label_path):
                self.label_map[idx] = label
                for file in os.listdir(label_path):
                    if file.endswith(".wav"):
```

```

        self.samples.append(os.path.join(label_path, file))
        self.labels.append(idx)

    def __len__(self):
        return len(self.samples)

    def __getitem__(self, idx):
        file_path = self.samples[idx]
        label = self.labels[idx]
        y, sr = librosa.load(file_path, sr=None)
        mfcc = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=13)
        feature = np.mean(mfcc.T, axis=0)
        return torch.tensor(feature, dtype=torch.float32), label

```

```

In [5]: train_dataset = AudioDataset("Data_People/Training") # type: ignore
        test_dataset = AudioDataset("Data_People/Testing") # type: ignore

        train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
        test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False)

```

Training

```

In [ ]: class AudioFNN(nn.Module):
        def __init__(self, input_size, num_classes):
            super(AudioFNN, self).__init__()
            self.model = nn.Sequential(
                nn.Linear(input_size, 64),
                nn.ReLU(),
                nn.Linear(64, 32),
                nn.ReLU(),
                nn.Linear(32, num_classes)
            )

        def forward(self, x):
            return self.model(x)

```

```

In [7]: input_size = 13
        num_classes = len(train_dataset.label_map)
        model = AudioFNN(input_size, num_classes) # type: ignore

        criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(model.parameters(), lr=0.003)

```

```

In [8]: # from torchsummary import summary

        # summary(model, input_size=(1, 28, 28))

```

```

In [9]: # Train in one epoch function
        def train_one_epoch(model, train_loader, loss_fn, optimizer, device):
            model.train()
            train_loss, train_correct = 0, 0

            for inputs, labels in train_loader:
                inputs, labels = inputs.to(device), labels.to(device)

                optimizer.zero_grad()
                outputs = model(inputs)
                loss = loss_fn(outputs, labels)
                loss.backward()
                optimizer.step()
                train_loss += loss.item() * inputs.size(0)
                _, predictions = torch.max(outputs, 1)
                train_correct += torch.sum(predictions == labels.data)

            return train_loss / len(train_loader.dataset), train_correct.double() / len(train_loader.dataset) # type: ignore

        # Validation function

```

```

def validate(model, val_loader, loss_fn, device):
    model.eval()
    val_loss, val_correct = 0, 0

    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            loss = loss_fn(outputs, labels)

            val_loss += loss.item() * inputs.size(0)
            _, predictions = torch.max(outputs, 1)
            val_correct += torch.sum(predictions == labels.data)

    return val_loss / len(val_loader.dataset), val_correct.double() / len(val_loader.dataset) # type: ignore

# Training and validation loop with timing
def train_and_validate(model, train_loader, val_loader, loss_fn, optimizer, epochs, device='cuda'):
    model.to(device)
    history = {
        'train_loss': [],
        'train_accuracy': [],
        'val_loss': [],
        'val_accuracy': []
    }

    for epoch in tqdm(range(epochs), desc="Training Progress", leave=True):
        epoch_start_time = time.time()

        train_loss, train_accuracy = train_one_epoch(model, train_loader, loss_fn, optimizer, device)
        val_loss, val_accuracy = validate(model, val_loader, loss_fn, device)

        history['train_loss'].append(train_loss)
        history['train_accuracy'].append(train_accuracy.item())
        history['val_loss'].append(val_loss)
        history['val_accuracy'].append(val_accuracy.item())

        epoch_end_time = time.time()

        tqdm.write(f'Epoch {epoch+1}/{epochs}: Train loss: {train_loss:.4f}, Train accuracy: {train_accuracy:.4f}, '
                  f'Val loss: {val_loss:.4f}, Val accuracy: {val_accuracy:.4f}, '
                  f'Time: {(epoch_end_time - epoch_start_time):.2f}s')

    return model, history

```

```

In [10]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
         model, history = train_and_validate(model, train_loader, test_loader, criterion, optimizer, epochs=30, device=device)

```

```

Training Progress:  3%|██████████| 1/30 [00:04<02:04, 4.30s/it]
Epoch 1/30: Train loss: 3.8544, Train accuracy: 0.4028, Val loss: 1.8342, Val accuracy: 0.3889, Time: 4.30s
Training Progress:  7%|██████████| 2/30 [00:05<01:01, 2.21s/it]
Epoch 2/30: Train loss: 2.3409, Train accuracy: 0.3750, Val loss: 1.2645, Val accuracy: 0.5556, Time: 0.74s
Training Progress: 10%|██████████| 3/30 [00:05<00:42, 1.56s/it]
Epoch 3/30: Train loss: 1.2278, Train accuracy: 0.5000, Val loss: 1.2531, Val accuracy: 0.3889, Time: 0.78s
Training Progress: 13%|██████████| 4/30 [00:06<00:31, 1.23s/it]
Epoch 4/30: Train loss: 1.1522, Train accuracy: 0.5833, Val loss: 1.1342, Val accuracy: 0.3889, Time: 0.72s
Training Progress: 17%|██████████| 5/30 [00:07<00:26, 1.07s/it]
Epoch 5/30: Train loss: 1.1086, Train accuracy: 0.4583, Val loss: 1.0132, Val accuracy: 0.5000, Time: 0.77s
Training Progress: 20%|██████████| 6/30 [00:08<00:22, 1.05it/s]
Epoch 6/30: Train loss: 0.7311, Train accuracy: 0.5278, Val loss: 0.8167, Val accuracy: 0.5556, Time: 0.72s
Training Progress: 23%|██████████| 7/30 [00:08<00:20, 1.13it/s]
Epoch 7/30: Train loss: 0.6713, Train accuracy: 0.6250, Val loss: 0.6978, Val accuracy: 0.6667, Time: 0.75s
Training Progress: 27%|██████████| 8/30 [00:09<00:18, 1.20it/s]
Epoch 8/30: Train loss: 0.5721, Train accuracy: 0.7500, Val loss: 0.6641, Val accuracy: 0.5556, Time: 0.73s
Training Progress: 30%|██████████| 9/30 [00:10<00:16, 1.24it/s]
Epoch 9/30: Train loss: 0.5004, Train accuracy: 0.7222, Val loss: 0.6668, Val accuracy: 0.6111, Time: 0.74s
Training Progress: 33%|██████████| 10/30 [00:11<00:15, 1.28it/s]

```

```

Epoch 10/30: Train loss: 0.4930, Train accuracy: 0.7639, Val loss: 0.5841, Val accuracy: 0.6667, Time: 0.73s
Training Progress: 37%|██████████| 11/30 [00:11<00:14, 1.28it/s]
Epoch 11/30: Train loss: 0.4633, Train accuracy: 0.7778, Val loss: 0.5869, Val accuracy: 0.6111, Time: 0.76s
Training Progress: 40%|██████████| 12/30 [00:12<00:13, 1.32it/s]
Epoch 12/30: Train loss: 0.4379, Train accuracy: 0.7639, Val loss: 0.5327, Val accuracy: 0.7222, Time: 0.71s
Training Progress: 43%|██████████| 13/30 [00:13<00:12, 1.34it/s]
Epoch 13/30: Train loss: 0.4283, Train accuracy: 0.7639, Val loss: 0.5070, Val accuracy: 0.7222, Time: 0.72s
Training Progress: 47%|██████████| 14/30 [00:13<00:11, 1.34it/s]
Epoch 14/30: Train loss: 0.3903, Train accuracy: 0.8194, Val loss: 0.5669, Val accuracy: 0.7222, Time: 0.74s
Training Progress: 50%|██████████| 15/30 [00:14<00:11, 1.35it/s]
Epoch 15/30: Train loss: 0.5028, Train accuracy: 0.7639, Val loss: 0.7538, Val accuracy: 0.6111, Time: 0.72s
Training Progress: 53%|██████████| 16/30 [00:15<00:10, 1.35it/s]
Epoch 16/30: Train loss: 0.4829, Train accuracy: 0.7639, Val loss: 1.0022, Val accuracy: 0.6111, Time: 0.74s
Training Progress: 57%|██████████| 17/30 [00:16<00:09, 1.36it/s]
Epoch 17/30: Train loss: 0.6829, Train accuracy: 0.6944, Val loss: 0.5724, Val accuracy: 0.6667, Time: 0.71s
Training Progress: 60%|██████████| 18/30 [00:16<00:08, 1.38it/s]
Epoch 18/30: Train loss: 0.4007, Train accuracy: 0.8056, Val loss: 0.5256, Val accuracy: 0.7778, Time: 0.71s
Training Progress: 63%|██████████| 19/30 [00:17<00:08, 1.36it/s]
Epoch 19/30: Train loss: 0.4365, Train accuracy: 0.8056, Val loss: 0.4090, Val accuracy: 0.8333, Time: 0.75s
Training Progress: 67%|██████████| 20/30 [00:18<00:07, 1.37it/s]
Epoch 20/30: Train loss: 0.3665, Train accuracy: 0.8333, Val loss: 0.4103, Val accuracy: 0.7778, Time: 0.72s
Training Progress: 70%|██████████| 21/30 [00:19<00:06, 1.36it/s]
Epoch 21/30: Train loss: 0.3431, Train accuracy: 0.8750, Val loss: 0.3899, Val accuracy: 0.8333, Time: 0.74s
Training Progress: 73%|██████████| 22/30 [00:19<00:05, 1.37it/s]
Epoch 22/30: Train loss: 0.3525, Train accuracy: 0.8194, Val loss: 0.4862, Val accuracy: 0.7778, Time: 0.72s
Training Progress: 77%|██████████| 23/30 [00:20<00:05, 1.37it/s]
Epoch 23/30: Train loss: 0.3260, Train accuracy: 0.8611, Val loss: 0.4300, Val accuracy: 0.8333, Time: 0.71s
Training Progress: 80%|██████████| 24/30 [00:21<00:04, 1.35it/s]
Epoch 24/30: Train loss: 0.3230, Train accuracy: 0.8889, Val loss: 0.6136, Val accuracy: 0.7222, Time: 0.76s
Training Progress: 83%|██████████| 25/30 [00:22<00:03, 1.37it/s]
Epoch 25/30: Train loss: 0.4578, Train accuracy: 0.7639, Val loss: 0.4184, Val accuracy: 0.8333, Time: 0.71s
Training Progress: 87%|██████████| 26/30 [00:22<00:02, 1.34it/s]
Epoch 26/30: Train loss: 0.3600, Train accuracy: 0.8194, Val loss: 0.3814, Val accuracy: 0.7778, Time: 0.77s
Training Progress: 90%|██████████| 27/30 [00:23<00:02, 1.35it/s]
Epoch 27/30: Train loss: 0.3637, Train accuracy: 0.8056, Val loss: 0.4532, Val accuracy: 0.8333, Time: 0.73s
Training Progress: 93%|██████████| 28/30 [00:24<00:01, 1.33it/s]
Epoch 28/30: Train loss: 0.2898, Train accuracy: 0.8611, Val loss: 0.4413, Val accuracy: 0.8333, Time: 0.77s
Training Progress: 97%|██████████| 29/30 [00:25<00:00, 1.35it/s]
Epoch 29/30: Train loss: 0.3231, Train accuracy: 0.8056, Val loss: 0.5257, Val accuracy: 0.7778, Time: 0.71s
Training Progress: 100%|██████████| 30/30 [00:25<00:00, 1.16it/s]
Epoch 30/30: Train loss: 0.3082, Train accuracy: 0.8472, Val loss: 0.3796, Val accuracy: 0.8333, Time: 0.77s

```

Result

```

In [14]: model.eval()

y_true, y_pred = [], []

with torch.no_grad():
    for features, labels in test_loader:
        features, labels = features.to(device), labels.to(device)
        outputs = model(features)
        _, predicted = torch.max(outputs.data, 1)
        y_true.extend(labels.cpu().tolist())
        y_pred.extend(predicted.cpu().tolist())

print('Accuracy of the model is:', accuracy_score(y_true, y_pred))

```


Accuracy of the model is: 0.8333333333333334

```

In [15]: # Confusion Matrix
cm = confusion_matrix(y_true, y_pred)
ConfusionMatrixDisplay(cm, display_labels=[train_dataset.label_map[i] for i in range(num_classes)]).plot()

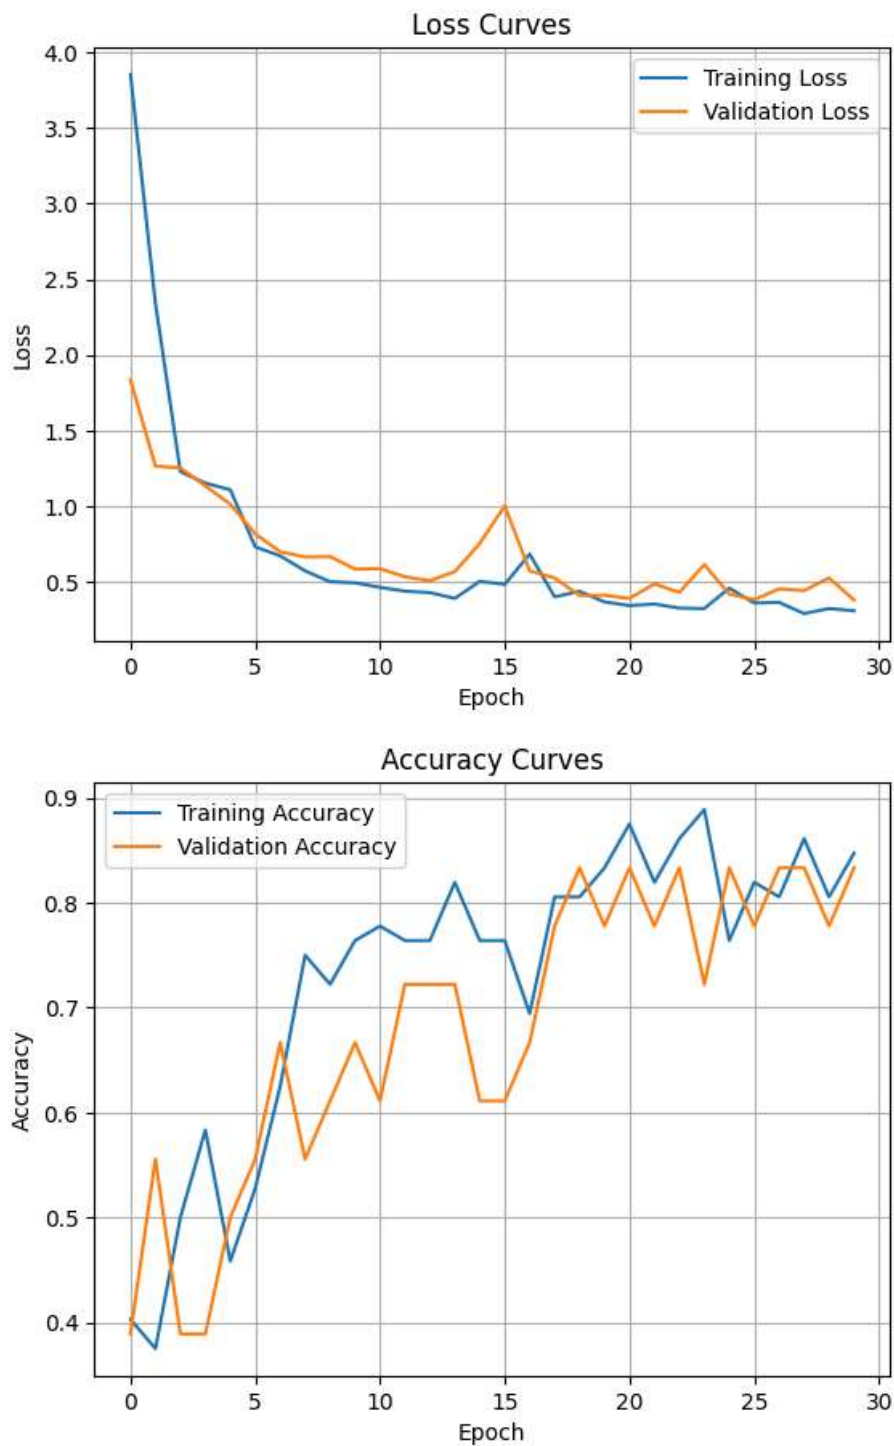
```

```
plt.title("Confusion Matrix")
plt.show()
```



```
In [13]: plt.figure()
plt.plot(history['train_loss'], label="Training Loss")
plt.plot(history['val_loss'], label="Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Loss Curves")
plt.legend()
plt.grid(True)
plt.show()

plt.figure()
plt.plot(history['train_accuracy'], label="Training Accuracy")
plt.plot(history['val_accuracy'], label="Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Accuracy Curves")
plt.legend()
plt.grid(True)
plt.show()
```



Discussion

This is my first model, a simple one using FNN. I also make a second model which is CNN which could be check after this and all of the other discussion in the CNN.